



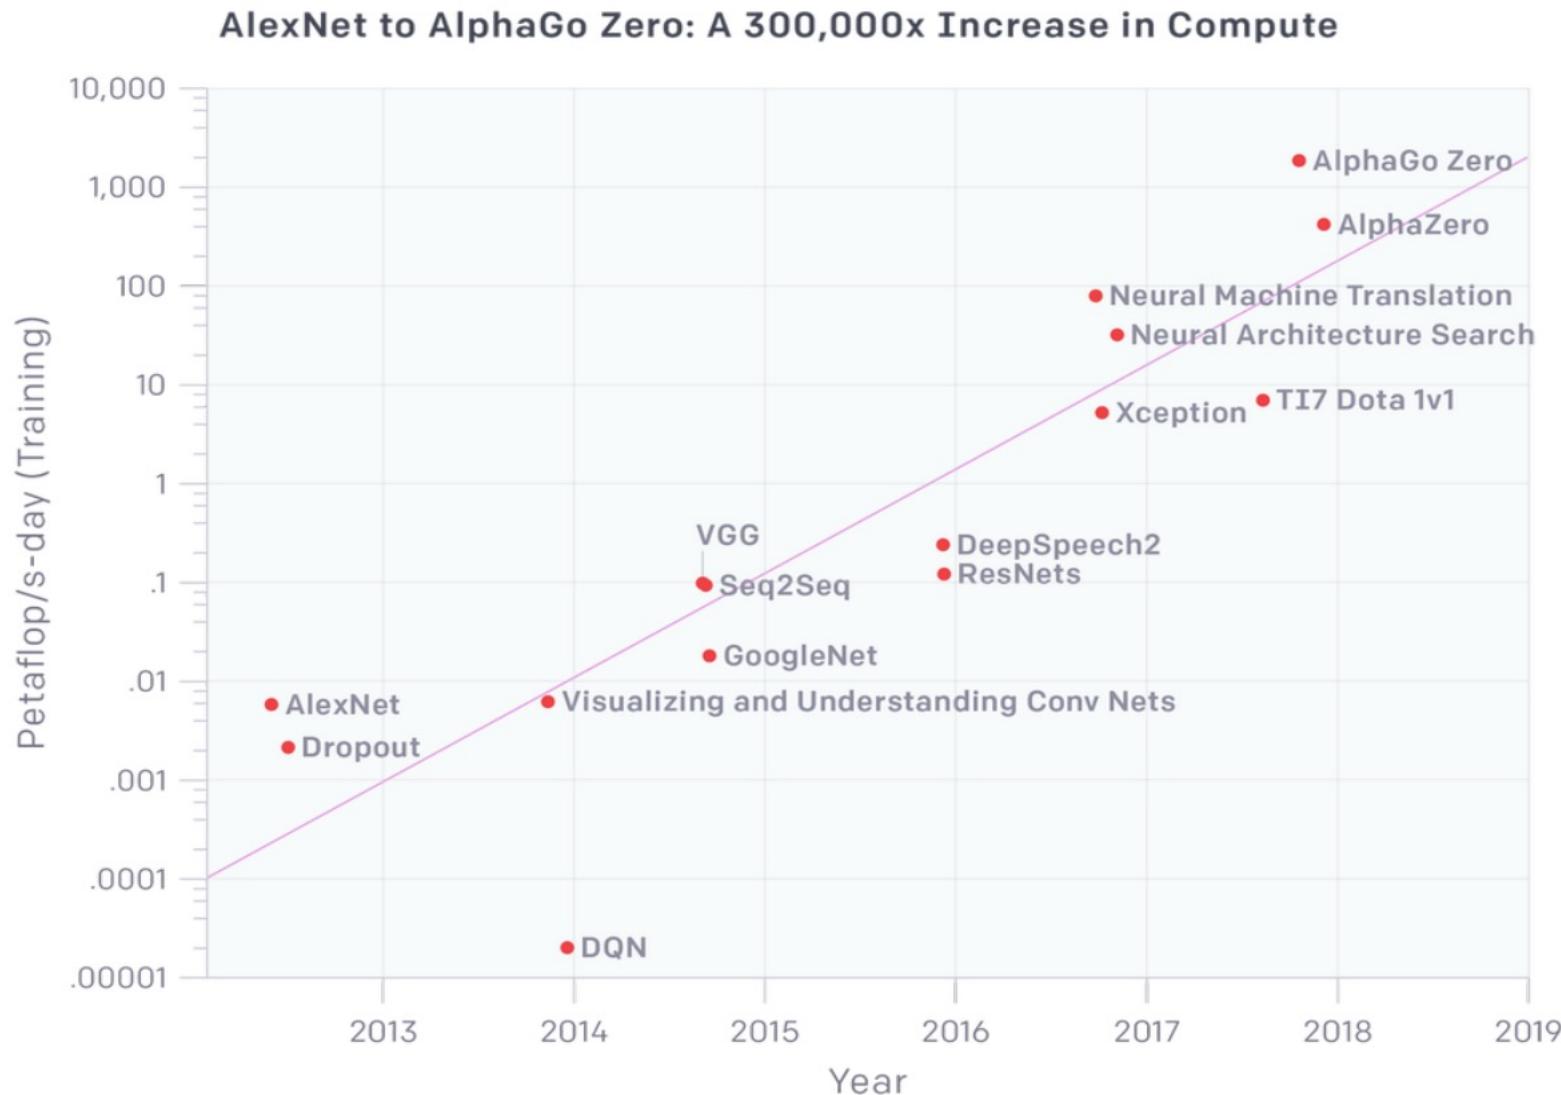
IT Infrastructure for Research: HPCs, Cloud Computing and beyond!

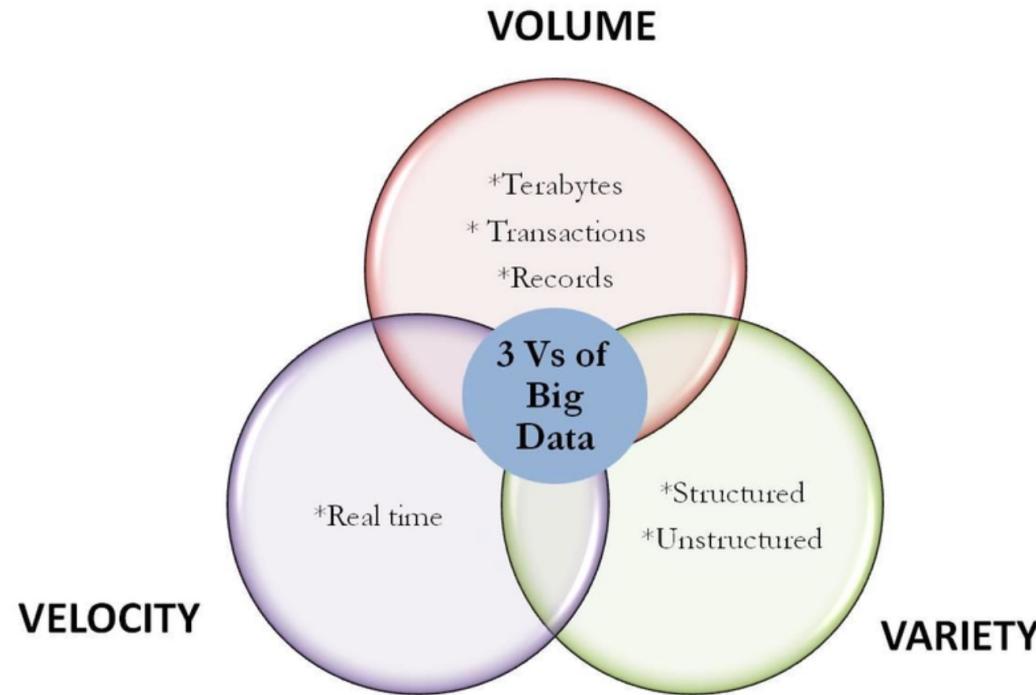
Mahsa Lotfi
April 19th , 2021

Special Thanks:
Mark Piercy
Riccardo Murri



Total Amount of Compute Power Used for different Machine Learning Models





We are exposed to “Big Data”
We need to be prepared!

A revolution for computing in science





Running hundreds of
tasks



Managing hundreds of
CPU hours for projects

It's too much workload management
We need to have tools!



Alpha Pattern:

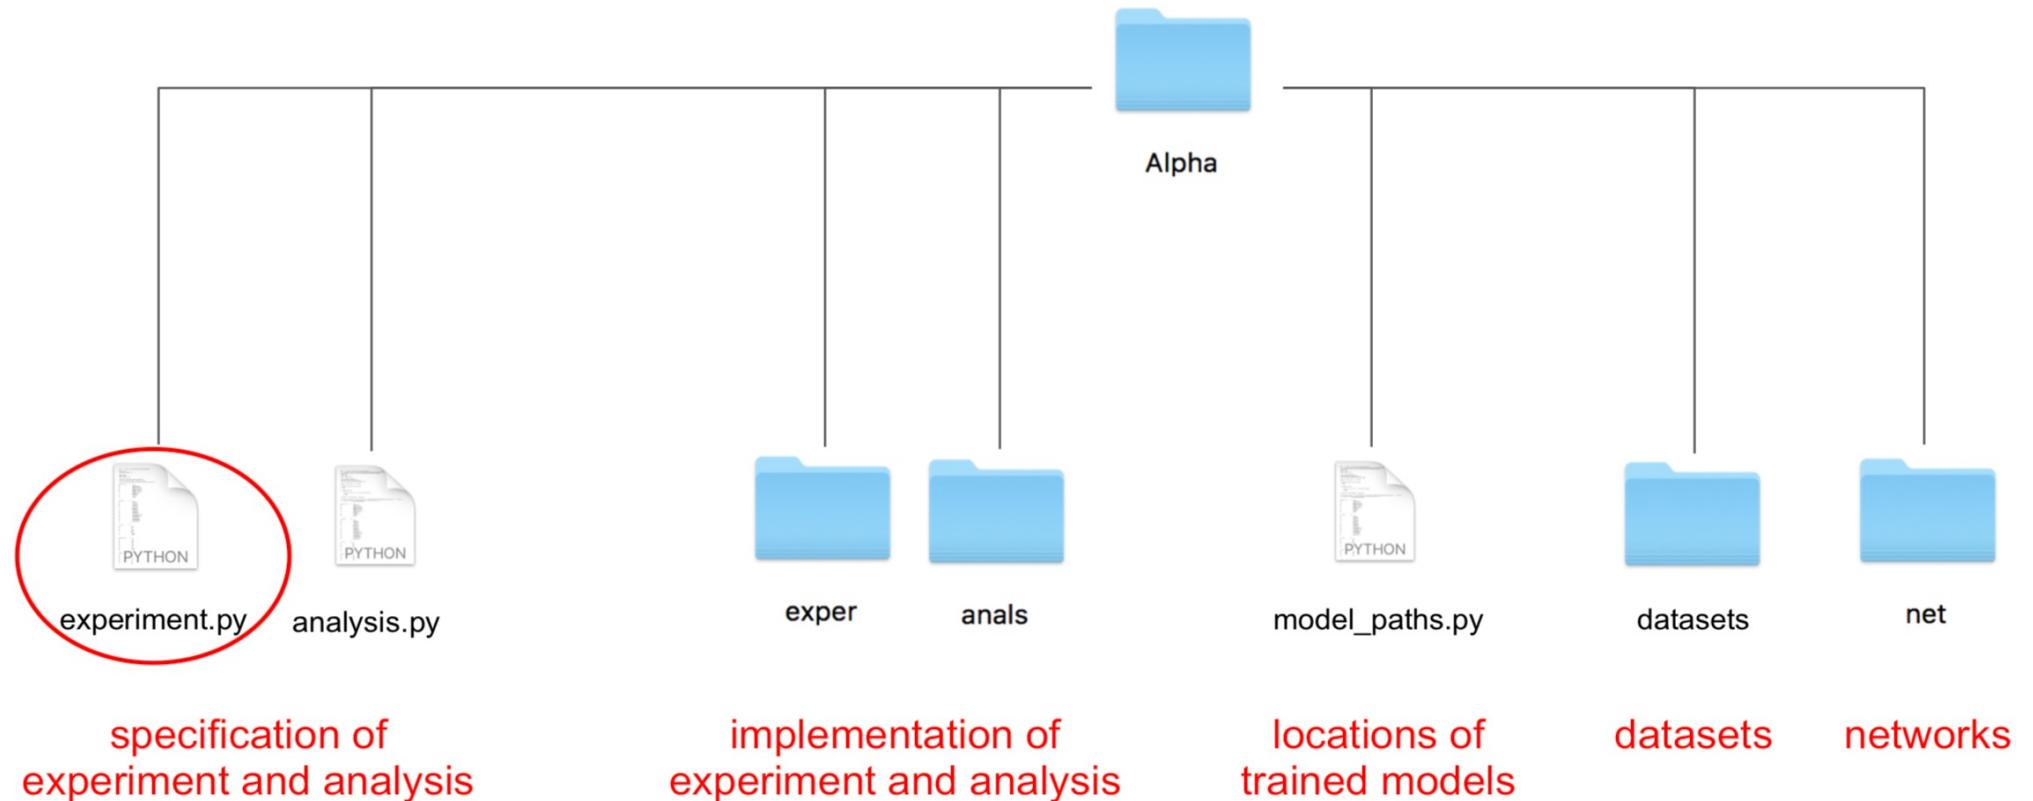
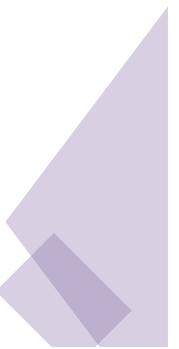
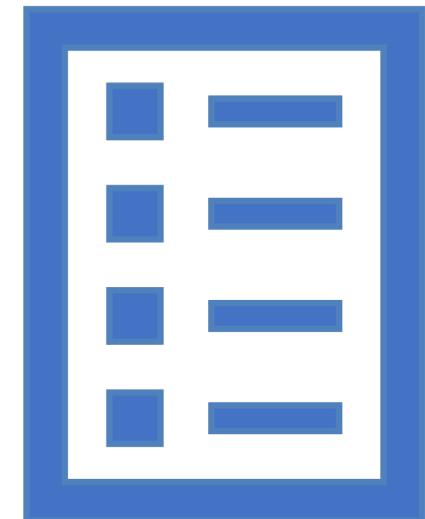


Table of Contents:

- Platforms for Scientific Computing
- High performance computing (HPC)
 - Clusters
 - Login Nodes
 - Compute Nodes
 - Why Using Clusters?
 - PCs vs. HPC
 - GPU
- SLURM
 - Fairshare
- Cloud Computing
- ElastiCluster
- ClusterJob
- ElastiCluster plus ClusterJob Demo

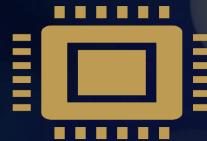


Platforms for Scientific Computing



Personal Workstations

Interactive use
Complete control over both the software and hardware
Limited computing power



Large shared batch-queuing systems

Centrally provided and administered
Usually have “Job Schedulers”
Standard “commodity servers” as compute nodes
High-performance network interconnecting nodes
Shared file system



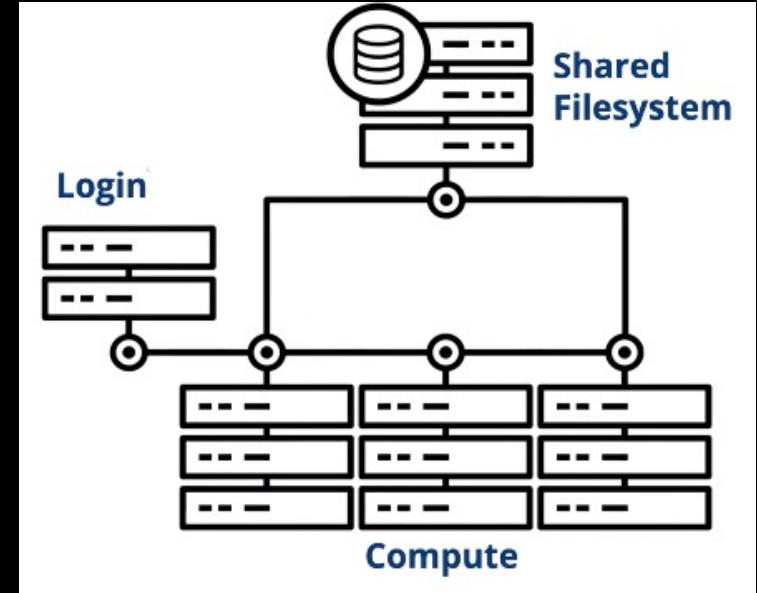
Cloud Computing

High Performance Computing (HPC)

- **HPC:** A collection of many separate servers/computers called “**nodes**”
- **Fast interconnections**
- **Distributed System:** components/nodes located on different networked computers
- **Parallel Computing:**
 - Large problems are divided into smaller ones
 - Simultaneously compute each smaller part

Cluster

- Multiple racks of computers without displays or keyboards
- Have common access to a common storage
- **Login Nodes:** Gateway to the cluster
- **Compute Nodes:** Computations are done here!
- **Shared Filesystem:** Presents data across all nodes



Login Nodes:

- Small number of head nodes (1 or 2)
- Access point for users to run jobs on the cluster
- Many users simultaneously log into the head node so no intensive jobs should run on the login node
- Good for basic tasks such as:
 - Uploading data
 - Managing files
 - Checking/ Managing jobs
 - Submitting jobs to the scheduler
 - Mostly has memory limit

Compute Nodes:

- Majority of computations running on these nodes
- Many more CPUs compared to regular computers (20-48)
- Big RAM compared to regular computers
- Some have accelerators (GPUs)
- GPU nodes: Both CPU cores and GPU available for running jobs

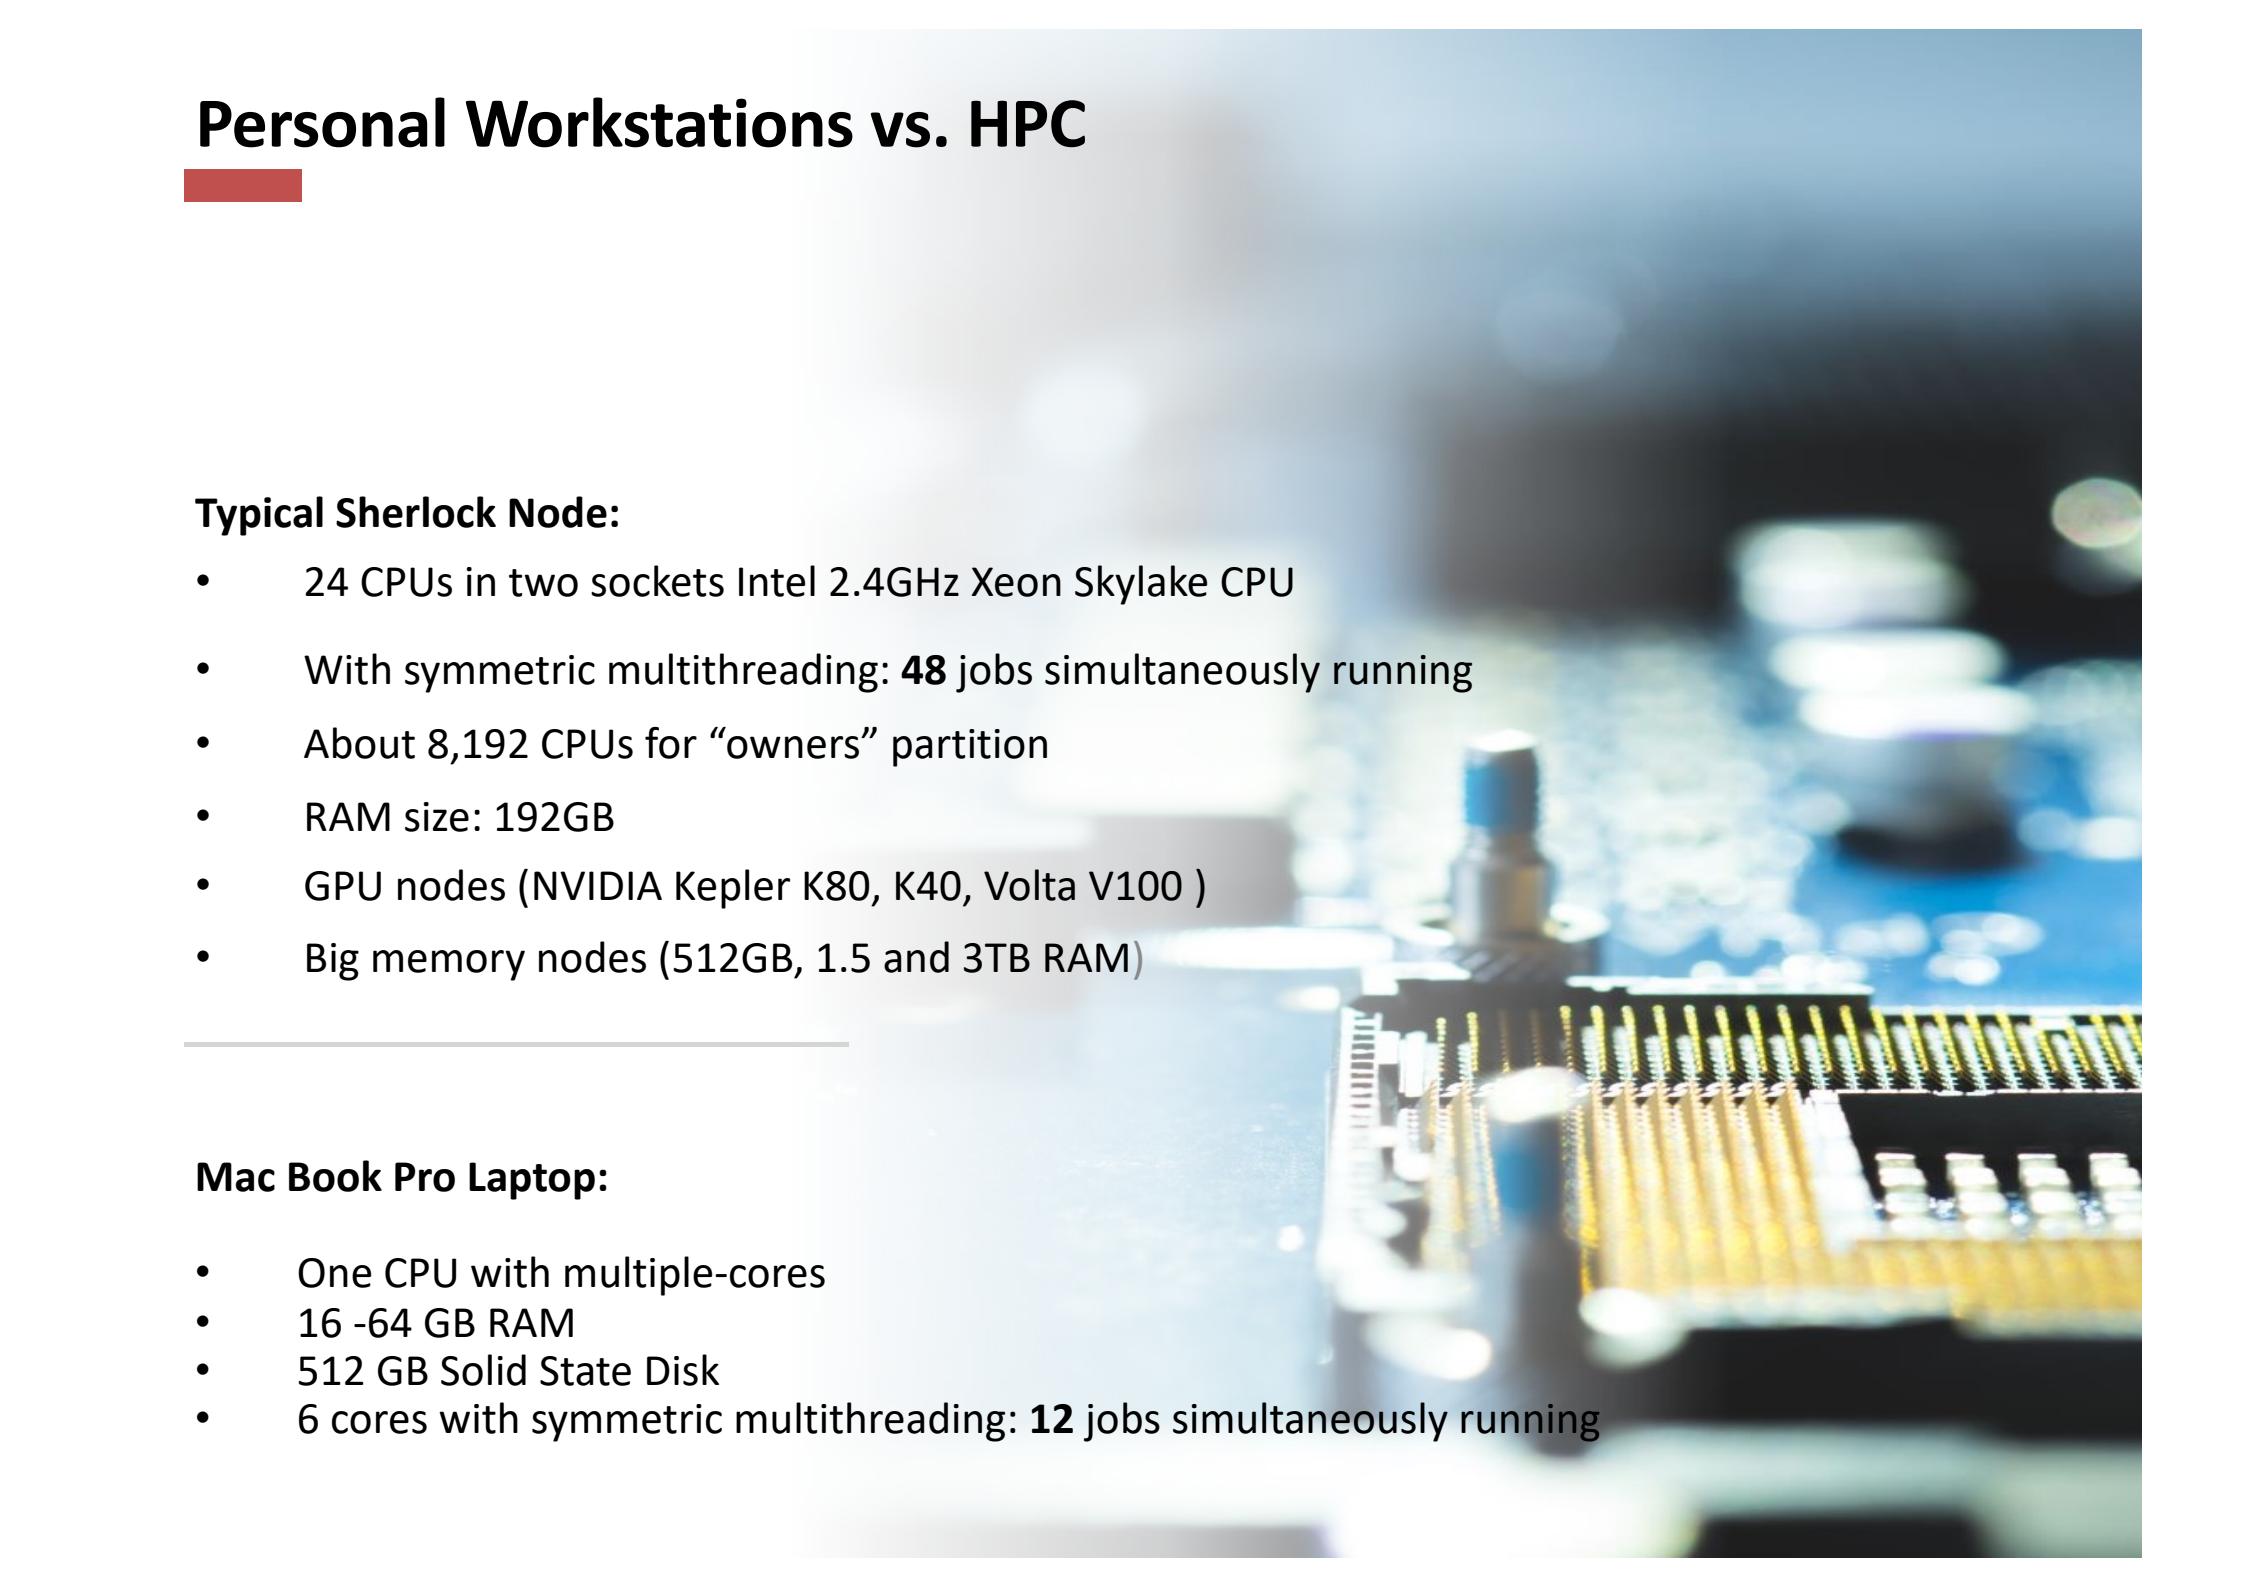
Why Using Clusters?

- ✓ Dealing with large datasets
- ✓ Need more disc space
- ✓ Need more memory
- ✓ Parallel jobs for faster results
- ✓ Using accelerators (GPUs)

Limitations:

- ✓ Not recommended if the jobs run for a long time (unless check-pointed)
- ✓ Not ideal for graphical tasks

Personal Workstations vs. HPC



Typical Sherlock Node:

- 24 CPUs in two sockets Intel 2.4GHz Xeon Skylake CPU
 - With symmetric multithreading: **48** jobs simultaneously running
 - About 8,192 CPUs for “owners” partition
 - RAM size: 192GB
 - GPU nodes (NVIDIA Kepler K80, K40, Volta V100)
 - Big memory nodes (512GB, 1.5 and 3TB RAM)
-

Mac Book Pro Laptop:

- One CPU with multiple-cores
- 16 -64 GB RAM
- 512 GB Solid State Disk
- 6 cores with symmetric multithreading: **12** jobs simultaneously running

C P U V E R S U S C O R E

CPU

An electronic circuit inside the computer that handles all instructions it receives from hardware and software running on the computer

A component inside the computer

A computer can have multiple CPUs or processors

CORE

Processing unit that receives instructions to carry on actions based on the instructions

Located inside the CPU

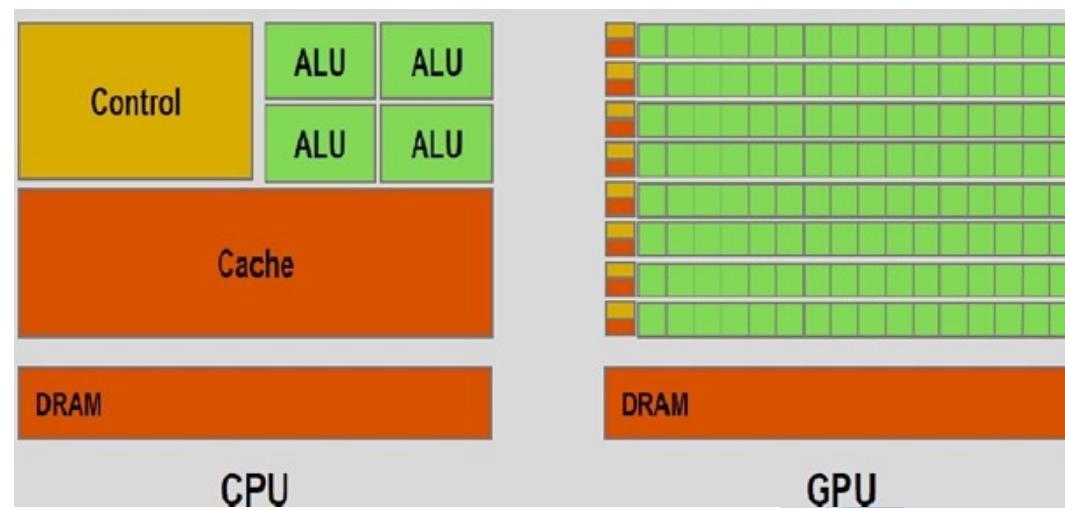
A CPU can have single or multiple cores

Visit www.PEDIAA.com

GPU: Graphic Processing Unit

- Highly parallel structure thus more efficient than general-purpose CPUs
- Good for parallel processing of large blocks of data
- Examples of GPU use cases:
 - Rendering graphics to a screen
 - Running Monte Carlo simulations
 - Multiplying large matrices for a machine learning algorithm
- More arithmetic logical units (ALUs) to calculate with

NOTE: Codes need to be modified to be GPU-compatible



Pros and Cons of HPC:

Pros:

- Servers are always on
- Accessible from anywhere by anyone in your PI group
- Much more compute power, hundreds of CPUs
- Large memory servers up to 3TB of RAM
- Job Scheduler

Cons:

- Learn how to use a job scheduler and the Linux shell
- Need permission for some software installations
- Wait in the queue for the “Job Scheduler”

What is SLURM?

Simple Linux Utility Resource Management



SLURM



- Open source, fault-tolerant and highly scalable cluster/workload management system
- Job scheduling system for large and small Linux clusters
- You need to tell the scheduler:
 - What resources you need such as # of CPUs, RAM, time, partition
 - Load the modules and run your code
 - Need to request as few resources as you need so your jobs pend for as small a time as possible
- [Why a management system?](#)
 - Managing and balancing the compute resources availability
 - Balancing the workloads

Fairshare

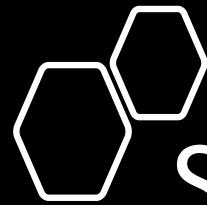


- Goal of having a [job queue](#):
 - Maximize utilization of the compute power
 - Ease the workload for users who do a lot of computation
 - Be fair to all users
- Used by SLURM to prioritize the tasks/jobs
- The more resources you use (CPU/RAM/Time/Nodes) in a 2 week sliding window, the lower your Fairshare score is and the more likely your jobs will wait in the queue

Fairshare (Cont.)



- Each job's priority in queue is determined by multiple factors, among them the user's Fairshare score
- Past usage computed based on a sliding window and progressively forgotten over time
- Stanford Sherlock uses *backfill*: smaller jobs can go in front of larger jobs, often regardless of the users Fairshare factor, thus increasing clusters utilization

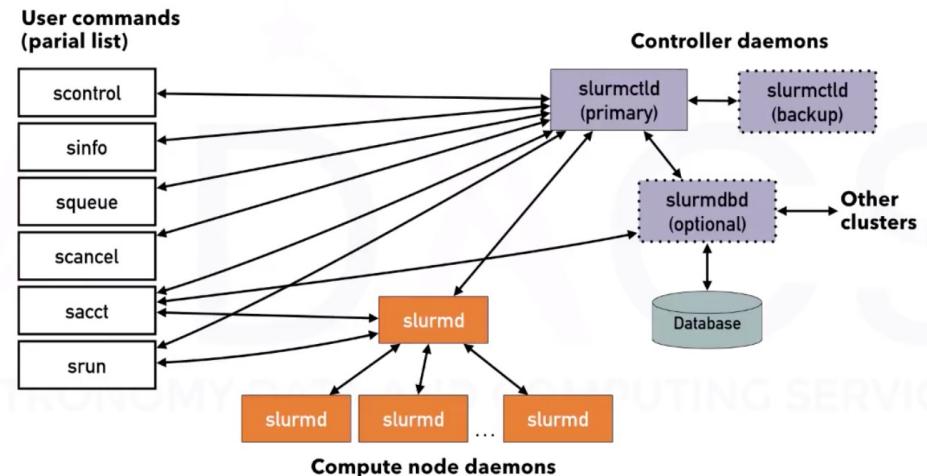


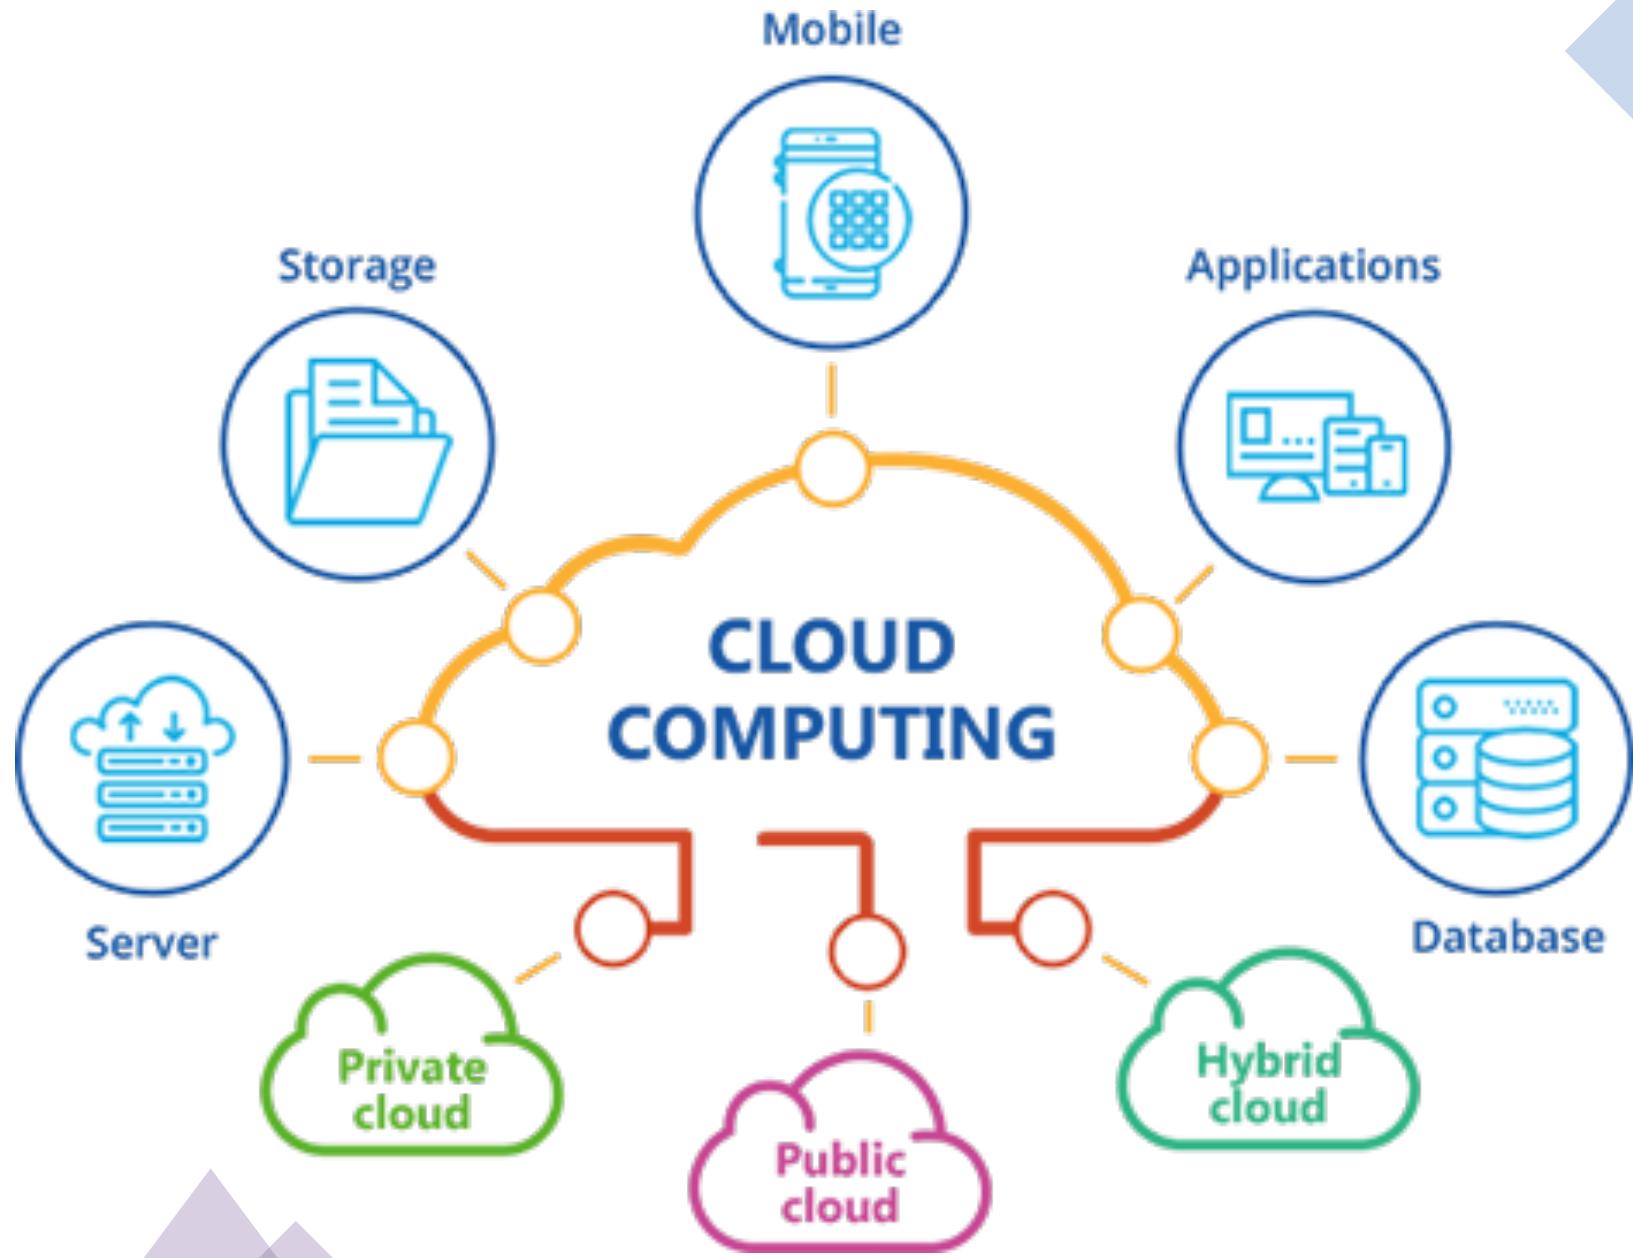
SLURM Architecture

- **Slurmd**: Slurm daemon running on each compute node
 - Accepting and monitoring tasks
 - Launching the tasks
 - Killing the tasks
- **Slurmctld**: Slurm central daemon running on management node
- Information queried by several commands:
 - Sacct
 - Salloc
 - Sattach
 - Sview
 - Sinfo
 - Scontrol
 - etc.



ARCHITECTURE



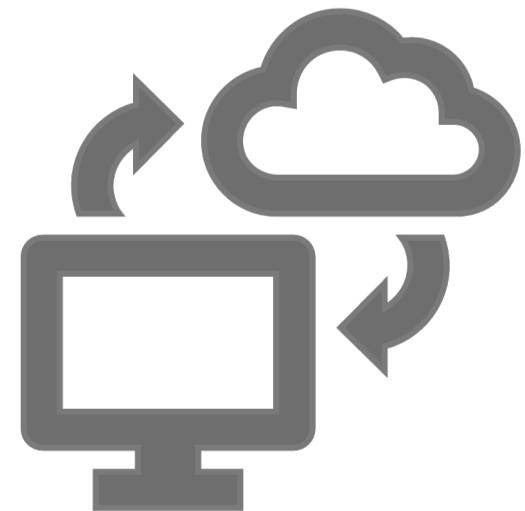


Embarrassingly Parallel Problems (EPP)

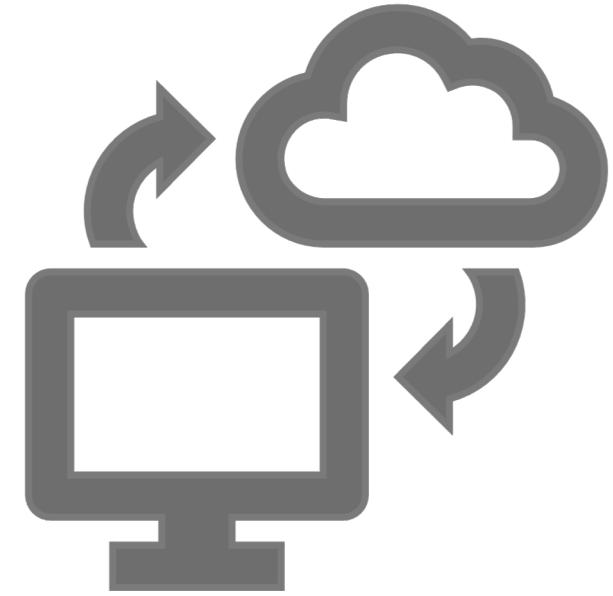
- Problems in which little/no effort is needed to separate the problem into a number of parallel tasks (Naturally Parallel Algorithms)
- The opposite of EPP: **Inherently Serial Problems** (cannot be parallelized at all)
- **Examples:**
 - DFT, each harmonic calculated independently
 - BLAST
 - Large Scale Face Recognition, etc.
- **Features:**
 - Almost no communication between the processes
 - Sub-solutions stored in disjoint memory locations
 - Sub-solution computations completely independent

Cloud Computing: IaaS

- **Infrastructure as a Service (IaaS)**
 - Virtualization of computing resources over the internet
 - Users log into the platform, create virtual machines, install OS etc.
 - **Examples:** AWS, GCE, Azure



Cloud Computing: PaaS



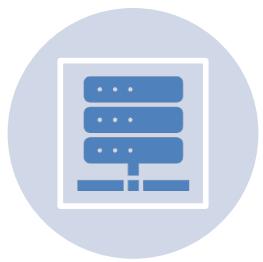
- **Platform as a Service (PaaS)**
 - Various hardware and software tools are available for application development to users over the internet
 - IT services available for users, accessible anywhere via a web browser
 - **Examples:** AWS Elastic Beanstalk, Google App Engine, Google Big Table, etc.

Cloud Computing: SaaS



- **Software as a Service (SaaS)**
 - Utilized for businesses in the cloud market
 - Utilizes the internet to deliver applications managed by a third-party vendor to its users
 - Directly through web browser, do not require any downloads or installations on the client side
- **Examples:** Google Workspace, Concur, Cisco WebEx, etc.
- **On-premises** software deployment: Software installed directly on the user's local machine, users have physical control over the hardware and the software

HPC vs. Cloud



- Good for EPP
- Good for large scale computing
- Mostly outperforms Cloud (Faster)
- Requires expensive hardware
- Fast interconnections
- Waiting for resources
- Learning a scheduler
- Not root, can't install some SWs



- Good for EPP
- Good for large scale computing
- Slow connections between nodes
- Cheaper than HPC
- Run on low cost commodity hardware
- No expensive HW/ SW upgrades
- No need to learn a job scheduler
- No waiting for resources not competing with hundreds of users for CPUs, RAM
- As a root ,install anything you want



ClusterJob

ClusterJob

- An automation system for high-throughput reproducible computations
- Easier parallelization of tasks
- CJ builds 'reproducible' computational packages that are easy to share with others
- Mainly written in Perl
- simple, easy-to-learn commands
- Currently supports MATLAB, Python and R
- Check pointing not necessary: Rerun the sub-problem

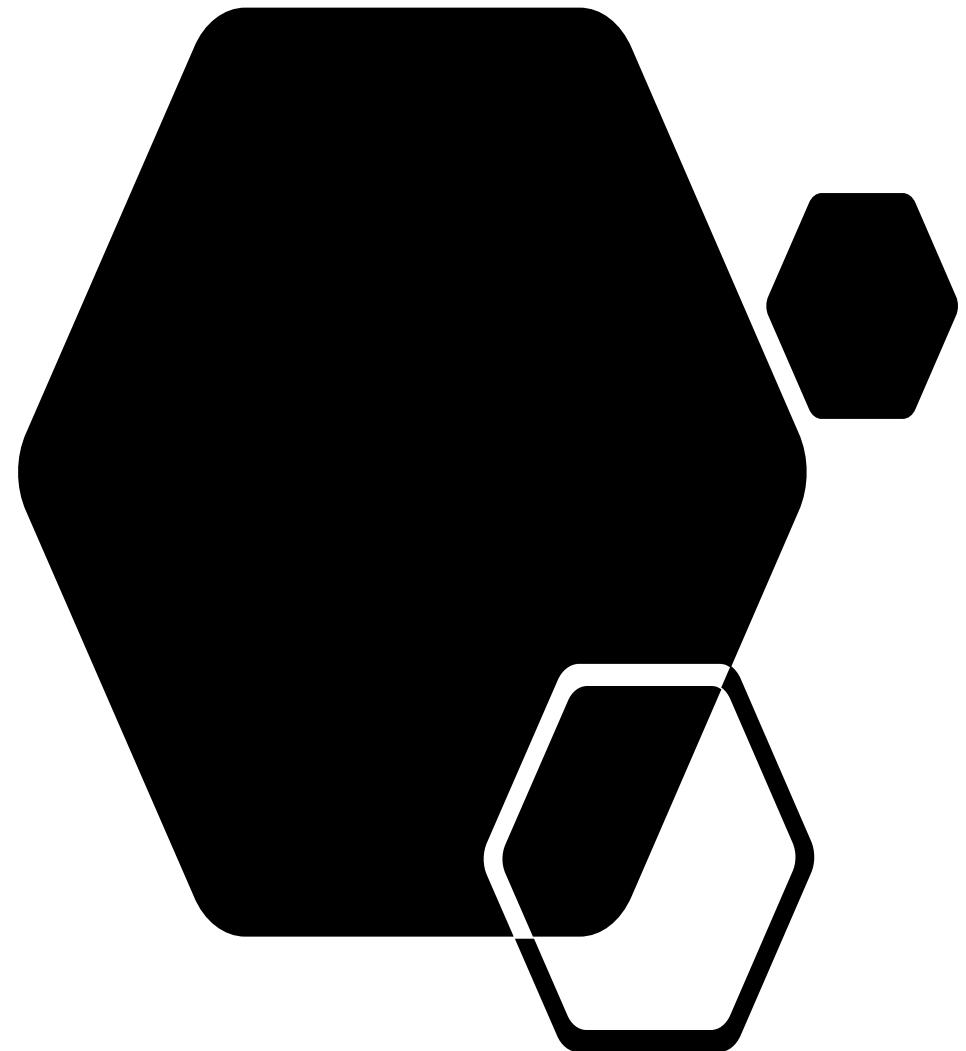


ClusterJob Useful Commands:



1. Write your Python/MATLAB/R code in a simple nested “for loop” format
2. Submitting jobs:
 - One job: `cj run file.py sherlock –dep Files –alloc “-p owners” –m “test”`
 - **Multiple jobs:** `cj parrun file.py sherlock –dep Files –alloc “-p owners” –m “test”`
3. Check the status of the jobs: `cj state PID`
4. Retrieve information: `cj log PID`
5. Gather all the results: `cj reduce PID`
6. Get the results in local machine: `cj get PID`

ElastiCluster



ElastiCluster

- Open-source software started at UZH
- Automated provisioning of virtual private clusters in the cloud
- Command line tool to **create, set up and scale** clusters with customized attributes and policies hosted on cloud
- Bespoke cluster up and running with a single command
- Additional commands can scale the cluster up and down

ElastiCluster Features:

- Supports several distros as base OS:
 - Debian
 - Ubuntu
 - CentOS
- Run on multiple clouds:
 - AWS
 - Google Cloud Engine
 - OpenStack
- **Issue: setup time grows linearly** with the number of cluster nodes

ElastiCluster Config File

```
# Create a cloud provider
[cloud/google]
provider=google
noauth_local_webserver=True
gce_client_id=*****
gce_client_secret=*****
gce_project_id=*****
zone=us-west1-b

[login/google]
image_user=ubuntu
image_user_sudo=root
image_sudo=True
user_key_name=elasticcluster
user_key_private=~/.ssh/id_rsa
user_key_public=~/.ssh/id_rsa.pub

[cluster/gce]
cloud=google
login=google
setup=ansible-slurm
security_group=default
frontend_nodes=1
compute_nodes=1
ssh_to=frontend
# Ask for 500G of disk
boot_disk_type=pd-standard
boot_disk_size=500

[cluster/gce/compute]
flavor=n1-standard-32
#flavor=n1-highmem-8
image_id=ubuntu-1604-xenial-v20171107b
accelerator_count=1
accelerator_type=nvidia-tesla-k80

[cluster/gce/frontend]
flavor=n1-standard-32
image_id=ubuntu-1604-xenial-v20171107b
```

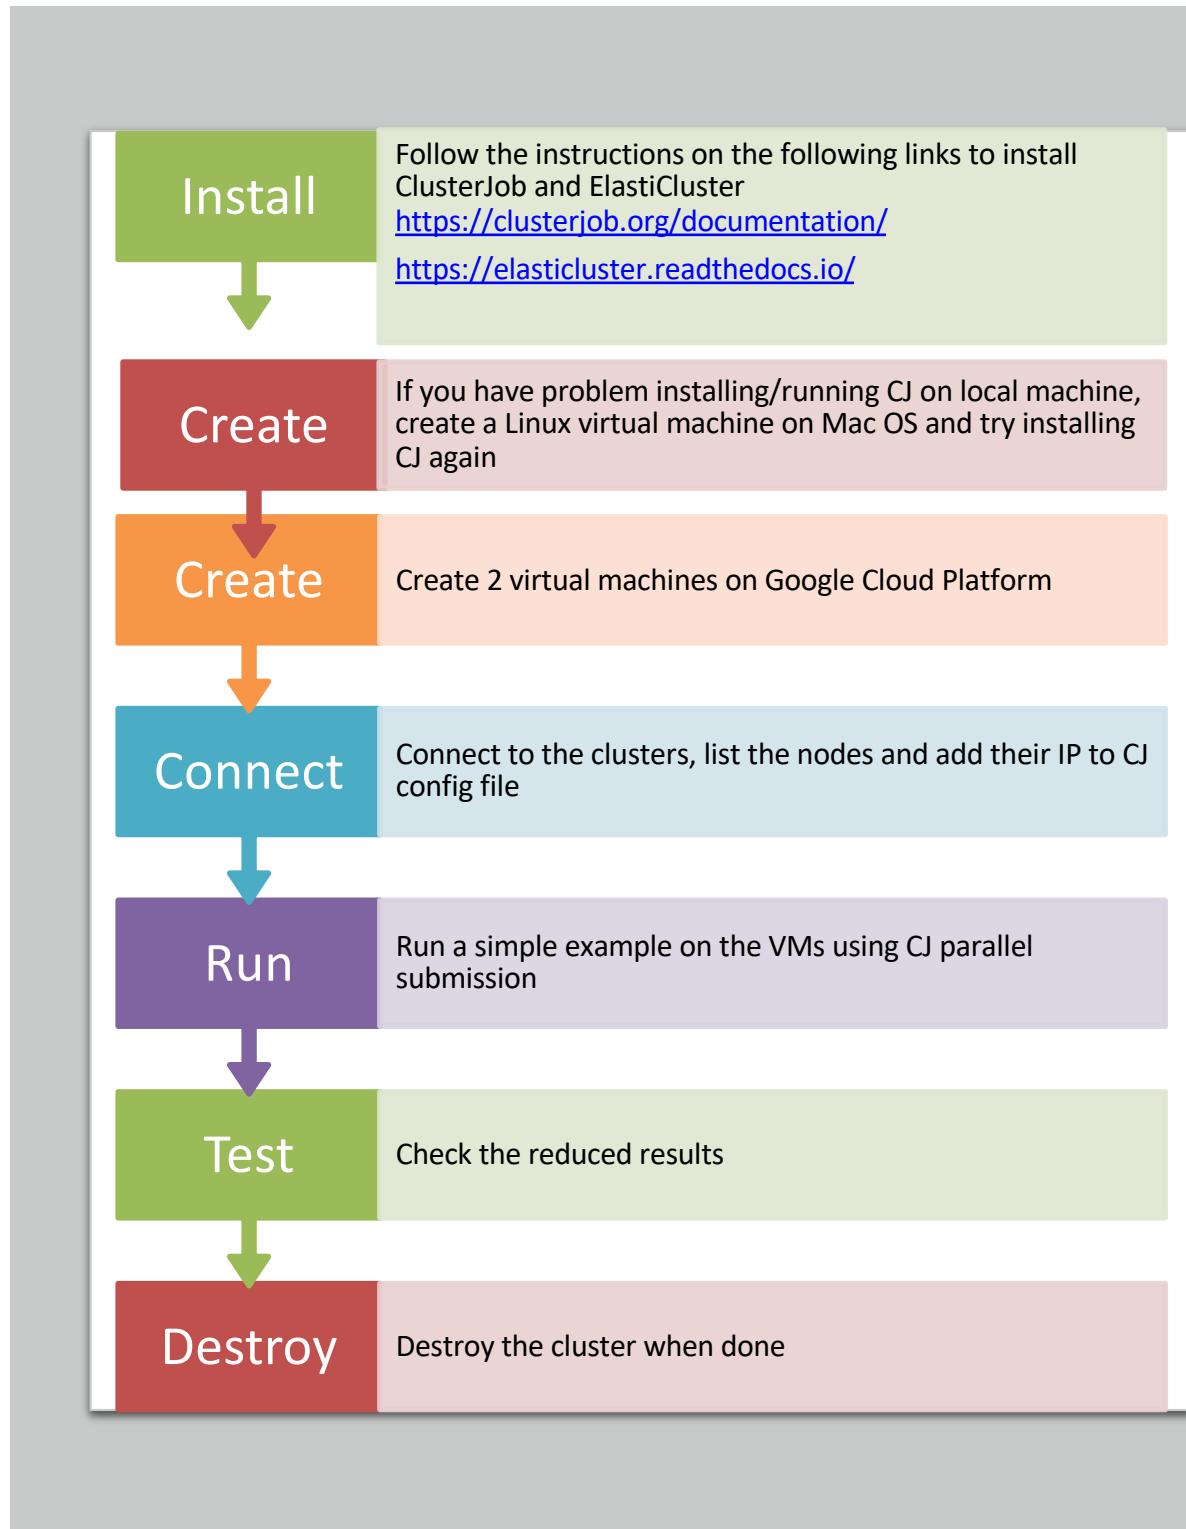
ElastiCluster-ClusterJob Model



+ ElastiCluster +



ElastiCluster + ClusterJob Demo



Issues regarding ElastiCluster?



- ElastiCluster source code:
 - <http://github.com/gc3-uzh-ch/elasticcluster> ElastiCluster
- Documentation:
 - <https://elasticcluster.readthedocs.org>
- Mailing-list:
 - elasticcluster@googlegroups.com
- Chat / IRC channel:
 - <http://gitter.im/elasticcluster/chat>



Riccardo Murri
University of Zurich, Switzerland

Acknowledgements



Mark Piercy
Stanford Research Computing Center



Thanks!
Any questions?